

Principles and roadmap
DevOps Cycle: plan->code->build->test->release->deploy->operate->monitor
Culture – Collaborative and customer-centred Automation – Integration testing, deployments, IaC Lean – Agile, scrappy, lean teams to minimize WIP (2-pizza rule) Measurement – Track and measure data to celebrate wins & pre-empt faults Sharing – Teach & learn from each other
Gene Kim's 3 ways: <ul style="list-style-type: none"> • Flow - Convert a business hypothesis into a technology-enabled service that delivers value to the customer • Learning – Ensure fast/constant feedback in all stages of the value stream • Feedback – Prioritize organizational learning and safety culture
Flow: <ul style="list-style-type: none"> • Make work visible • Limit WIP • Reduce batch size and handoffs • Identify and elevate constraints • Eliminate waste in the value stream
Feedback <ul style="list-style-type: none"> • Monitoring to see problems as they occur • Swarming problems to build new knowledge • Quality at source • Optimize for downstream work centres
Continual learning and experimentation <ul style="list-style-type: none"> • Improve daily work • Local discoveries to global improvements • Resilience patterns that introduce chaos and tension • Leadership – creating conditions for success
Microservice: Anything that provides functionality over a network – logically represents a business capability, otherwise treated as a black box that can be independently developed/deployed
Microservice architecture: An assembly of fine-grained services that allow independent, continuous deployment
Docker
Choreography: Services communicate with each other asynchronously
Orchestration: Composite services invoke atomic services to fulfil business functions
Atomic service: Provides functionality related to 1 capability. Self-contained and do not depend on any other services
Drawbacks of microservices: <ul style="list-style-type: none"> • Microservices make things worse, not better, if developers have a poor testing culture • Decomposing a service incorrectly gives a distributed monolith • Distributed systems are inherently more complex
Docker networking: All traffic is routed through docker engine. Containers can communicate with each other within the network, but outside applications must use forwarded ports as specified by Dockerfile/docker-compose.yml
Continuous Integration
Test pyramid (descending order of difficulty, brittleness and cost): <ul style="list-style-type: none"> -UI testing -End-to-end testing -Component testing

<ul style="list-style-type: none"> -Integration testing -Unit testing -Code analysis
Continuous Deployment
Continuous delivery – releases are automated, deployments are manual Continuous deployment – deployments are fully automated
Benefits: Lower failure rates, faster feedback, faster flow, reliable releases Tradeoffs: Building binaries for diff platforms, customer perceptions of CD, only works with a testing culture
Release versioning: Traditionally semantic versioning <major>.<minor>.<patch> <ul style="list-style-type: none"> • Major: breaking, incompatible API changes • Minor: backwards-compatible functionalities • Patch: backwards compatible bug fixes
Serverless deployment: Abstracts away infra & resources, invoking microservices implemented as lambda functions. <ul style="list-style-type: none"> • Pros: Very lightweight and scalable • Cons: Cold starts, large functions, not suitable for long-running jobs
Service as a container: Package as a Docker image and deploy each service as a container. Each service has its own IP and file system <ul style="list-style-type: none"> • Pros: Portable, isolated, constrained resources, encapsulation • Cons: Responsibility for administering images and specifying/administering container infrastructure
AWC ECS concepts: <ul style="list-style-type: none"> • Task definition: Describes containers that form your app – images and resource constraints • Service: Runs and maintains a specified number of tasks (can set auto-scaling as well) • Cluster: Logical grouping of tasks and services. Able to orchestrate containers across multiple EC2 instances
Deployment patterns <ul style="list-style-type: none"> • Rolling: Gradually replace container instances, specifying a minimum healthy % • Blue-green: Transfer all traffic to new container instances; old instances on 'standby' for rollback • Canary: Release to a small subset of users first
Feature flags: Allow new features to be enabled/disabled on toggle, promoting easy rollback, graceful degrades, resilience in deploys
Microservices communication
Styles: 1-1 sync(HTTP, gRPC), 1-1 async(fire & forget, async request/reply), 1-many async (pub/sub)
Message-oriented middleware: Acts as a broker for any communication style
Communication patterns <ol style="list-style-type: none"> 1. Orchestration – Stateless composite service that manages atomic services 2. Choreography – Atomic services communicate with each other asynchronously 3. Choreo with process engine – Combines async behaviour & process visibility
AMQP: A public message queueing protocol that runs on top of TCP <u>Publisher</u> -> <u>Exchange</u> -> <u>Queue</u> -> <u>Subscriber</u> . Publishers publish message to an exchange, exchanges are bound to queues using a routing key (direct/fanout/topic), subscribers consume messages from queues.
Saga pattern: Create a set of compensating transactions for every local transaction that a microsvc makes. Since microservices are not isolated, concurrency is still an issue.
Solutions: RabbitMQ (message broker for AMQP), prog. language implementations (pika, etc.)
API gateways

Strangler fig pattern: Gradually decouple the system from top-down, starting with edge services and working to core. Transform -> Co-exist -> Eliminate

API gateway patterns:

1. Composition – aggregating requests to atomic svcs
2. Protocol translation
3. Edge functions (auth, rate limiting, caching, metric, logging), consider *separation of concerns*

+ Encapsulation/façade
+ Reduces network overhead
+ Simplify client-side interactions
+ Loose coupling

- More management
- Must be highly available
- May cause development bottleneck
- No clear ownership

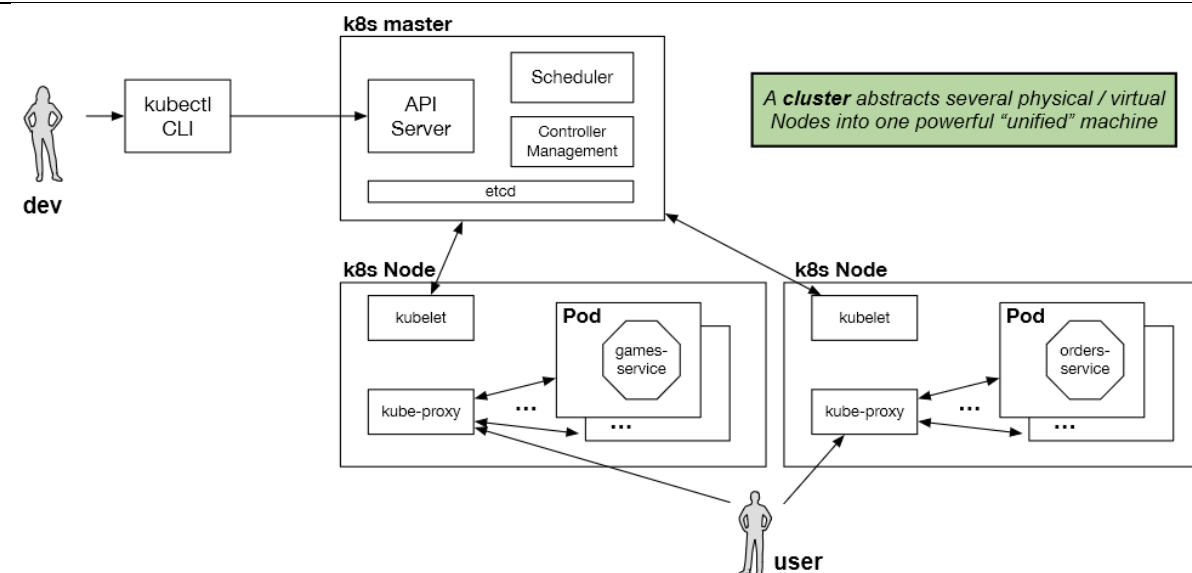
General purpose: One API gateway for all frontend clients
- General purpose, multiple responsibilities, no clear ownership (against DevOps style)

Backends for frontends (BFF): Each frontend (mobile/web/native) has its own gateway
- Frontend oriented, more separated responsibilities, easier tracking, possible code duplication

Solutions: AWS API gateway, Kong, DIY (nginx, zuul, etc.)

Kubernetes

Declarative model: Use manifest files (YAML) to specify desired state, and k8s takes actions to achieve that desired state



Node: Contains services required to run pods

Kubelet: Agent that runs on a node, ensuring that pods are running and healthy

kube-proxy: Allows communication over virtual network

Pod: A basic unit of deployment, with at least 1 container sharing an internal IP address and storage volume. Meant to be ephemeral and replaced

API-server: Allows dev to interact with k8s deployment through UI, CLI or REST calls

Controller manager: Keeps track of cluster state and makes adjustments if necessary

Scheduler: Watches for newly created pods and selects nodes for them to run on

Service: Abstraction that provides a static network location that exposes some pods. Internal addresses of Pods may change, but service ip/port does not change

- *ClusterIP:* Service reachable on an internal IP
- *NodePort:* Service reachable on Node's IP at a static port, externally reachable at <NodeIP>:<NodePort>
- *LoadBalancer:* Service exposed externally using cloud provider's load balancer (1-to-1)

Ingress: Exposes multiple Services through a single cloud load balancer

Ingress controller: Implementations include cloud providers, nginx, etc

Rolling updates: Default rolling update has a limit of 25% unavailable pods/extra pods Auto-scaling rules can be set based on CPU usage, min or max replicas
Monitoring
Telemetry: Automated process of collecting/transmitting metrics from remote points to monitoring systems. 1-Instrument systems, 2-Collect metrics from systems, 3-Monitor through alerts and visualizations
Challenges: Number of components, data quantity, silos, maintenance burden
What to monitor?: <i>Work metrics</i> – Throughput, success, errors, latency <i>Resource metrics</i> – Utilization, saturation, availability <i>Events</i> - Code changes, alerts, scaling events <i>CI/CD</i> – Time taken for pipeline stages, pass/fail events
Log levels: DEBUG -> INFO -> WARN -> ERROR -> FATAL, severity level dictates incident response
Self-service: Metrics should be easily accessible, “info radiators” w/o privileged access
Analysis: Through statistical operations or anomaly detection (independent of probability distribution)
Prometheus: Centralized telemetry software, stores metrics in time series database, taken from service endpoints in a PULL model. Query language PromQL allows selection and aggregation of data. Data can then be visualized through software like <i>Grafana</i>
Infrastructure as Code
IAC tools: <ul style="list-style-type: none"> • Shell scripts • Configuration management: Install/manage software on servers • Provisioning: Create ephemeral resources (servers, brokers, security groups)
Terraform <ul style="list-style-type: none"> • <code>init</code>: Initialize current working directory and prepare files for use w Terraform • <code>fmt</code>: Rewrite .tf files to a canonical style • <code>plan</code>: Reads remote state, updates local state, compares local state with .tf files and proposes changes to make remote object match desired configuration • <code>show</code>: Displays local state snapshot • <code>destroy</code>: Destroys remote infra managed by Terraform
Importing infrastructure <ul style="list-style-type: none"> • Declare a blank resource type and resource name • <code>terraform import resource_type.resource_name <remote.id></code> • Copy local state into configuration file, making sure to delete any ephemeral resource fields (e.g. ARN, ID, IP address)